

Experiments

Tim Visée & Nathan Bakhuizen

October 2018



Intro

In this document, we will list all the motions that are recognized by *Can't Touch This*. These experiments can be conducted by first setting up the platform itself. Instructions for this can be found in the manual. After setting everything up, conducting the experiments is a piece of cake. Start up the platform, open a webbrowser and go to [.](#) Click on 'Visualize' to view your fingers being traced.

Please try to make the following gestures above the LeapMotion sensor. You will receive visual feedback if a gesture has been made succesfully.

Accuracy & Reliability

- Straight line
- Circle clockwise
- Circle counter-clockwise
- Big circle clockwise
- Big circle counter-clockwise
- Triangle clockwise
- Triangle counter-clockwise
- Mini square clockwise
- Square clockwise
- Square counter-clockwise

Gesture variety

Visual feedback

Recognition consistency

It is important that all gestures are successfully performed and can be recognized multiple times. This is because recognizing gestures can be error-prone. By performing gestures often, the risk of it recognizing a gesture incorrectly decreases dramatically.

Detection speed

Another important aspect of the gesture recognition system is the speed, or rather, the delay for detecting a performed gesture. To avoid confusion, we've decided to define the term delay as follows: *The time it takes for the platform to show the name of the correct gesture in the console, after a gesture has been completely performed above the sensor.*

Counting frames

Since it can be tricky to accurately measure the time difference between an invocation and an action, we've decided to record 20 invocations. With the resulting footage, we'll be able to count delay in frames, with precision accurately enough for our use case. The gesture that we choose was the square, as it has quite distinctive corners.

Results

After performing the our own tests it was quite clear that delay varies greatly. The *delay* varied from about 5 to 11 frames. With a speed of 60 frames per second, each frame is about 17 milliseconds. This defines a delay varying between 85 to 187 milliseconds. There were two outliers with a delay of 16 to 22 frames which translate to about 270 and 370 milliseconds.

We believe the varying difference has to do with the accuracy of detecting the gesture itself. Imagine you are drawing a perfect square gesture (perfect relative to the square gesture you might have previously recorded) above the sensor, the gesture will be recognized quicker before even fully completing the gesture, as compared to a square motion that differs slightly. Because it's virtually impossible to perform the gesture in exactly the same way every iteration, a difference in detection is indeed expected.

Note that our method of testing latency is still quite simplistic. It doesn't take computer and screen delay into account. Nor does it test what the latency of the used sensor itself is. Our test strictly focusses on the delay between performing the gesture and seeing visual feedback. We performed the test with the visualizer enabled, which was visible on the material we recorded. What's interesting is that it looks like the gesture is instantly (meaning; within the same frame) recognized as the last sampled point of a gesture shows up on the visualizer. This would suggest that the actual latency for detection on the processed data is faster than 17 milliseconds. This is of course not very scientific, but we feel it's an awesome result non the less.

What is important though, is how responsive the detection feels. During these tests, the detection (strictly speaking about detection speed) felt snappy even during the case of those outliers. The cool thing is that after you've performed a gesture you don't have to wait for a detection notification before the next gesture can be performed, that helps quite a lot for it to feel responsive to our opinion when when repeatedly experimenting with the sensor.

The platform doesn't support binding generic actions to gestures to control a computer yet. And until that is tested, it's hard to say whether we'd feel the same when controlling a computer with these gestures. We mention this idea of binding actions later in this report in the *Ideas* section.

Ideas

During the project, and during the experiments we've conducted we came up with quite a few ideas. Some of them are things we'd wanted to implement but didn't have the time for it due to time constraints. Other ideas originate from test results we've collected. Here follow a few of them:

Multiple fingers

The current system only has support for gestures consisting of a single trace, being the trace drawn by the index finger of a hand. The first logical expansion would of course be to support gestures having motions with 5 fingers. Sadly, due to time constraints, we couldn't implement this.

Our guess is that implementing this isn't too hard when using the same detection system as our current platform. The implementation might be as easy as extending our `Hand` structure (in code) to contain 5 traces, instead of just 1. This might be cool to experiment with for other groups taking a look at this project.

Not just finger tips

With the extension described above, the system would still only be tracking finger tips. The sensor we're currently using is able to report the position of all joints in a hand, including the center of each bone and the center of the hand palm. This is easily visible in the 3D visualizer that is provided with the LeapMotion software package. The `leap-rs` wrapper that is currently used doesn't have an interface to obtain this information yet, but this could be implemented without too much effort.

The same might be true for this implementation, just tracking the trace for every joint might be enough. We're of course not sure, that's something cool to test out.

3D detection

Although we intended to support it at first, 3D gesture detection would be an awesome improvement. The current system uses 2D rotational data to represent user motions. Simply said; because of processing data to this format, the third dimension is lost. We found it difficult to reliably transform a 3D curve into 2D rotation data, and wanted to get the 2D system working first.

Of course, motions in 3D space are transformed into data our platform uses. During this process the depth (Z-axis) is lost though. Drawing a circle sideways would be detected as moving up and down on a straight line. In this system, you might imagine drawing on a virtual plane on the X and Y axis. Supporting the Z-axis as well for a 3rd dimension would open up a lot of new possibilities for drawing gestures in different orientations and would support much more complex gestures.

Bindable actions

Our idea was to allow users to bind custom actions to recorded gesture templates quite early in the project. This would allow a user to configure the platform to control their computer in a generic way. The implementation can be as simple as binding mouse (and scroll wheel) actions, keys and system commands.

Such an implementation would immediately make the project much more viable for use by others (given that it would work properly). This also opens new research questions such as; can gestures be complex enough to support enough gestures for controlling a generic computer?

This idea was also the segway to research on how usable a gesture detection system is for use in medical & sterile environments. Imagine doctors not having to touch a röntgen machine to swipe between pictures, by using touch-less hand gestures instead. The implementation of bindable actions should make quite clear whether such a system is usable and reliable.

Multiple sensors

An awesome extension would be to use and combine multiple sensors in the platform, instead of using a single LeapMotion sensor. The obvious choice would be to use a second LeapMotion sensor to combine data of the two. This could be used to filter jitter, and to fallback if the view for one sensor is occluded.

This is one of the first things we looked into during the project. Sadly, it became apparent that the LeapMotion SDK doesn't support connecting multiple sensors, even though the provided library would suggest otherwise as there are functions available to obtain a list of connected sensors. [This](#) forum thread is about this topic, in which users are asking for support to use multiple sensors. In this thread, LeapMotion developers have said this is not possible.

It would be fun to use a different computer, or possibly a virtual machine with a virtual USB controller to support connecting another sensor. The data from this sensor would then need to be transmitted to the platform to collect the sensor data. This is however outside the scope of our project, and is a challenge to properly implement on its own.

Another option is to implement support for other kinds of sensors, such as the Xbox Kinect, or something different. Our platform would then need to provide the proper abstractions to aggregate the data from different kinds of sensors into something the platform could use.