Experiments

Tim Visée & Nathan Bakhuijzen

October 2018

Preface

In this document, we will go through some experiments we've conducted for the platform we have built during the research minor KB-80. For each experiment we've written a new section describing what it is about, how we conduct the experiment and what our findings are. The general idea of conducing these experiments is to assert whether the current platform is *usable* for the purpose we had originally in mind, namely controlling a generic computer.

You may of course conduct similar experiments as we've described along with reading this document. Make sure you read the *Manual* on how to install and use the product.

The last part of the document goes over some ideas we came up with when working on on the project, and when conducting these experiments. These ideas may be implemented by others that continue the work on this project.

Accuracy & Reliability

- Straight line
- Circle clockwise
- Circle counter-clockwise
- Big circle clockwise
- Big circle counter-clockwise
- Triangle clockwise
- Triangle counter-clockwise
- Mini square clockwise
- Square clockwise
- Square counter-clockwise

Gesture variety

Visual feedback

An interesting aspect for using our platform is visual feedback, and what the effect of using it is on users using our platform. With visual feedback, a user would be able to see on a computer screen how the computer (or rather; sensor) recognizes his hand(s) and/or gestures. The general idea for this test is that showing visual feedback might improve accuracy, as users could adapt the gesture they're performing on what they see the computer recognizes. The sensor isn't super accurate all of the time, or users might go outside the view of the sensor. The visual feedback would immediately show this.

Visualizers

Two visualizers are available. The first one is part of the LeapMotion software package, and is a 3D visualizer that can be started from the LeapMotion control panel. This visualizer shows a 3D representation of recognized hands above the LeapMotion sensor inside the bounding box of the sensor. Joints are visible in the hand, although our software only uses the position of the index finger tip at this moment.

The second visualizer is the one we've implemented. It is show on the web interface that can be opened after starting the software. This visualizer is quite different, not only in aesthetics, but also in what it shows. Our visualizer is 2D and shows the data that has been processed by our software. The *Manual* explains how we are processing the data, and that we're mapping 3D points from the sensor in traces of rotational data. This is what is rendered in our visualizer, and is essentially what is used for gesture detection. This is therefore *closest to reality*, where *reality* consists of user defined gestures in our software and the recognition system. If gestures show up correctly in this visualizer, the software should recognize them correctly. This same visualizer is used for recording and trimming new gestures.

Results

For this test we've tried a lot of different gestures, in a lot of different positions. Here are some of our statements on our findings:

- Drawing circles isn't a problem. The only thing a visualizer might help for is to see whether you're performing the gestures within the view of the sensor, and to see whether your hand is properly recognized.

- Without visual feedback attempting to draw squares, can easily result in drawing rectangles, especially when focussing on the screen instead of looking at your hand. Having the 2D web visualizer helps greatly as you can see whether you're drawing a square properly.

- The 3D visualizer is easier to understand, although the data shown in the 2D visualizer is more consistent. This, in the sense that using the 2D visualizer to guide you to perfectly draw a gesture works better and produces more accurate results than using the 3D visualizer.

- Either visualizer is useful to see whether the sensor is jittering, causing detection failures.

- The 2D web visualizer is useful to figure out why a false positive may have been detected as it clearly shows the shape that is used for recognition.

Based on these statements we can conclude that using a visualizer is a good idea, as it does seem to improve results for most gestures. There are gestures however that don't really benefit from using a visualizer. Such gestures can be used in a situation where a visualizer isn't preferred as it would use up screen real estate. The accuracy of a user drawing a gesture correctly might improve the more the user uses the system. We haven't tested this.

None of the visualizers notify the user if a drawn gesture is similar to a user defined gesture (but is not a match). Imagine that you're drawing a circle, but too large for it to be detected. The visualizers don't show that your circle is too large. This might be something to improve on.

Note that we, the developers of the platform, have conducted this test. Other users that don't understand the internals of the system might experience this different. The second visualizer in the web interface shows processed data, this might look weird and unintuitive for other users.

Recognition consistency

It is important that all gestures are successfully performed and can be recognized multiple times. This is because recognizing gestures can be error-prone. By performing gestures often, the risk of it recognizing a gesture incorrectly decreases dramatically.

Detection speed

Another important aspect of the gesture recognition system is the speed, or rather, the delay for detecting a performed gesture. To avoid confusion, we've decided to define the term delay as follows: *The time it takes for the platform to show the name of the correct gesture in the console, after a gesture has been completely performed above the sensor*.

Counting frames

Since it can be tricky to accurately measure the time difference between an invocation and an action, we've decided to record 20 invocations. With the resulting footage, we'll be able to count delay in frames, with precision accurately enough for our use case. The gesture that we choose was the square, as is has quite distinctive corners.

Results

After performing the our own tests it was quite clear that delay varies greatly. The *delay* varied from about 5 to 11 frames. With a speed of 60 frames per second, each frame is about 17 milliseconds. This defines a delay varying between 85 to 187 milliseconds. There were two outliers with a delay of 16 to 22 frames which translate to about 270 and 370 milliseconds.

We believe the varying difference has to do with the accuracy of detecting the gesture itself. Imagine you are drawing a perfect square gesture (perfect relative to the square gesture you might have previously recorded) above the sensor, the gesture will be recognized quicker before even fully completing the gesture, as compared to a square motion that differs slightly. Because it's virtually impossible to perform the gesture in exactly the same way every iteration, a difference in detection is indeed expected.

Note that our method of testing latency is still quite simplistic. It doesn't take computer and screen delay into account. Nor does it test what the latency of the used sensor itself is. Our test strictly focusses on the delay between performing the gesture and seeing visual feedback. We performed the test with the visualizer enabled, which was visible on the material we recorded. What's interesting is that it looks like the gesture is instantly (meaning; within the same frame) recognized as the last sampled point of a gesture shows up on the visualizer. This would suggest that the actual latency for detection on the processed data is faster than 17 milliseconds. This is of course not very scientific, but we feel it's an awesome result non the less.

What is important though, is how responsive the detection feels. During these tests, the detection (strictly speaking about detection speed) felt snappy even during the case of those outliers. The cool thing is that after you've performed a gesture you don't have to wait for a detection notification before the next gesture can be performed, that helps quite a lot for it to feel responsive to our opinion when when repeatedly experimenting with the sensor.

The platform doesn't support binding generic actions to gestures to control a computer yet. And until that is tested, it's hard to say whether we'd feel the same when controlling a computer with these gestures. We mention this idea of binding actions later in this report in the *Ideas* section.

Ideas

During the project, and during the experiments we've conducted we came up with quite a few ideas. Some of them are things we'd wanted to implement but didn't have the time for it due to time constraints. Other ideas originate from test results we've collected. Here follow a few of them:

Multiple fingers

The current system only has support for gestures consisting of a single trace, being the trace drawn by the index finger of a hand. The first logical expansion would of course be to support gestures having motions with 5 fingers. Sadly, due to time constraints, we couldn't implement this.

Our guess is that implementing this isn't too hard when using the same detection system as our current platform. The implementation might be as easy as extending our `Hand` structure (in code) to contain 5 traces, instead of just 1. This might be cool to experiment with for other groups taking a look at this project.

Not just finger tips

With the extension described above, the system would still only be tracking finger tips. The sensor we're currently using is able to report the position of all joints in a hand, including the center of each bone and the center of the hand palm. This is easily visible in the 3D visualizer that is provided with the LeapMotion software package. The `leap-rs` wrapper that is currently used doesn't have an interface to obtain this information yet, but this could be implemented without too much effort.

The same might be true for this implementation, just tracking the trace for every joint might be enough. We're of course not sure, that's something cool to test out.

3D detection

Although we intended to support it at first, 3D gesture detection would be an awesome improvement. The current system uses 2D rotational data to represent user motions. Simply said; because of processing data to this format, the third dimension is lost. We found it difficult to reliably transform a 3D curve into 2D rotation data, and wanted to get the 2D system working first.

Of course, motions in 3D space are transformed into data our platform uses. During this process the depth (Z-axis) is lost though. Drawing a circle sideways would be detected as moving up and down on a straight line. In this system, you might imagine drawing on a virtual plane on the X and Y axis. Supporting the Z-axis as well for a 3rd dimension would open up a lot of new possibilities for drawing gestures in different orientations and would support much more complex gestures.

Bindable actions

Our idea was to allow users to bind custom actions to recorded gesture templates quite early in the project. This would allow a user to configure the platform to control their computer in a generic way. The implementation can be as simple as binding mouse (and scroll weel) actions, keys and system commands.

Such an implementation would immediately make the project much more viable for use by others (given that it would work properly). This also opens new research questions such as; can gestures be complex enough to support enough gestures for controlling a generic computer?

This idea was also the segway to research on how usable a gesture detection system is for use in medical & sterile environments. Imagine doctors not having to touch a röntgen machine to swipe between pictures, by using touchless hand gestures instead. The implementation of bindable actions should make quite clear whether such a system is usable and reliable.

Multiple sensors

An awesome extension would be to use and combine multiple sensors in the platform, instead of using a single LeapMotion sensor. The obvious choice would be to use a second LeapMotion sensor to combine data of the two. This could be used to filter jitter, and to fallback if the view for one sensor os occluded.

This is one of the first things we looked into during the project. Sadly, it became apparent that the LeapMotion SDK doesn't support connecting multiple sensors, even though the provided library would suggest otherwise as there are functions available to obtain a list of connected sensors. This forum thread is about this topic, in which users are asking for support to use multiple sensors. In this thread, LeapMotion developers have said this is not possible.

It would be fun to use a different computer, or possible a virtual machine with a virtual USB controller to support connecting another sensor. The data from this sensor would then need to be transmitted to the platform to collect the sensor data. This is however outside the scope of our project, and is a challenge to properly implement on it's own.

Another option is to implement support for other kinds of sensors, such as the XBox Kinect, or something different. Our platform would than need to provide the proper abstractions to aggregate the data from different kinds of sensors into something the platform could use.